Architecting Provably Fair XR Art: A Two-Tiered Framework of Invariant Envelopes and Vector-Ranked Governance

Foundational Architecture: Controlled-Invariant Envelope Geometry

The foundational element of the proposed research framework is the development of controlled-invariant envelope geometry, which serves as the primary mechanism for guaranteeing minimum standards of safety, ecological responsibility, and economic viability for all participants in smart-city creative experiences . This concept, rooted in viability theory, establishes a multi-dimensional "envelope" around each stakeholder's state vector, representing their protected status across various critical axes . The core principle is that no project plan, algorithmic process, or deployment strategy can be approved if it would drive any stakeholder outside their guaranteed envelope . This provides a non-regressive guarantee, ensuring that optimization efforts never come at the cost of fundamental rights or baseline well-being . The user has explicitly prioritized this component for immediate implementation, recognizing its necessity for bounding physical and ecological parameters in public-facing XR deployments before higher-level governance structures can be meaningfully applied .

The mathematical formalism for these envelopes must be "spreadsheet-grade," meaning every derived metric must be an explicit algebraic formula or finite-sum rule directly computable from ALN (Artificial Life Network) telemetry . This requirement elevates the concept from an abstract theoretical construct to a deployable engineering specification. The implementation strategy involves encoding each fairness kernel, K, as a polytope—a geometric object with flat sides—which mirrors the existing approach used for other viability kernels like the `bio.safety.envelope.citizen.v1` . This choice of representation is crucial because polytopes allow for efficient computation of distance and feasibility. From this polytope-based representation, two critical metrics can be derived as closed-form spreadsheet cells: "kernel distance" and "exit-risk" . The kernel distance quantifies how close a given project trajectory is to violating the stakeholder's minimum guarantees. It could be formally defined as the minimum distance from the current state vector to the boundary of the polytope K. For a polytope defined by a set of linear inequalities $A\xi \leq b$, where ξ is the state vector, the distance can be computed using standard methods from computational geometry. The second metric, "exit-risk," goes a step further by modeling the probability that disturbances—such as market shifts, policy changes, or unforeseen events—could push the system trajectory outside the kernel . This aligns with the principles of Control Barrier Functions (CBF), where a function $h_i(\xi)$ measures the distance to a constraint boundary, and its time derivative is controlled to remain positive, ensuring safety . By calculating this risk in real-time from ALN data, the system can proactively reject or auto-correct plans that exhibit a high likelihood of failure, thus enforcing reversibility and safety .

This framework generalizes beyond socio-economic coordinates to include bio-mechanical, cognitive, legal, and ecological axes, creating a universal calculus for invariance that can be applied to any cybernetic or socio-technical system . For the immediate goal of smart-city music and art projects, these axes translate into tangible, enforceable constraints on the XR-grid environment itself. For instance, the "haptics" axis would bound the intensity of physical feedback to prevent discomfort or harm, while the "sound intensity" axis would cap decibel levels to protect auditory health and avoid disturbing the surrounding urban fabric. The "crowd density" axis would define maximum occupancy limits for virtual stages or plazas to ensure public safety and accessibility, and the "energy use" axis would constrain the power consumption of XR nodes to align with broader sustainability goals . These physical parameters are not arbitrary; they are direct consequences of the underlying stakeholder fairness vectors, which are themselves bounded by the controlled-invariant envelopes. An artist's creative ambition to create a powerful sensory experience is therefore channeled within a safe and reversible operational space defined by these multi-axis viability kernels. The system proves its own fairness by design: any plan that attempts to increase system utility (e.g., artistic impact or revenue) by pushing a stakeholder below their guaranteed floor is mathematically rejected during the viability check .

The mathematical research tasks for this priority area involve proving the conditions under which a specified polytope K remains invariant under the combined effects of bounded disturbances w and control inputs u for different system classes, such as XR grids or smart-city flows . This requires deriving explicit formulas for the kernel distance and exit-risk functionals that satisfy monotonicity and convexity properties, ensuring that the controller cannot inadvertently cross a soul.guardrail envelope . Furthermore, these metrics must be compositional, meaning that when subsystems (e.g., individual XR nodes or project teams) are combined into a larger system (e.g., a city-wide festival), the invariance property is preserved under clearly stated conditions . This composability is essential for analyzing large-scale, complex deployments without them becoming opaque black boxes . The creation of Rust crates to implement these functions is a critical next step, providing the low-level computational engine needed to make these abstract mathematical objects operational within Cybercore-Brain . These crates would handle the encoding of fairness kernels as polytopes, compute the associated distance and exit-risk metrics for every candidate plan, and feed these values into the decision-making pipeline . This transforms the framework from a static model into a dynamic, responsive control law that augments human creators and operators, allowing them to experiment and innovate while being shielded from systemic risks .

Metric	Mathematical Representation (Conceptual)	Purpose & Computation
Fairness Kernel (K) / Viability Kernel	Polytope in $\mathbb{R}^n$: $K = \xi \in \mathbb{R}^n \| A\xi \leq b$	Defines the hard lower bounds for a stakeholder's state vector (pay, voice, credit, ecology). Computed from ALN rows.
Kernel Distance	$d(\xi, K) = \min_{\xi' \in \partial K} \|\xi - \xi'\|$	Measures the minimum slack over active inequalities defining the polytope. A closed-form spreadsheet cell indicating proximity to a violation.
Exit-Risk Function (CBF-style)	$\rho(\xi) = \min_{i \in I(\xi)} \frac{c_i^T \dot{\xi}}{\|c_i\|}$	Computes the rate of change of the distance to the boundary. If negative, the trajectory is exiting the kernel. Used for proactive rejection.

In essence, the controlled-invariant envelope geometry provides the immutable foundation upon which the entire fair and safe ecosystem is built. It acts as a digital constitution, written in the language of mathematics, that defines the non-negotiable minimum terms of participation for every entity involved in a smart-city creative project. Its immediate implementation is paramount because it creates the necessary conditions for trust and responsible innovation. Without these provable, automated guardrails, any discussion of higher-order fairness concepts like Pareto optimality becomes moot. By making this architecture directly computable and integrable with existing systems, the framework moves decisively from theoretical possibility to practical reality, ready for deployment in the challenging environments of smart cities.

Operationalizing Stakeholder Protections Across Four Dimensions

The core objective of the research framework is to provide concrete, role-specific protections for stakeholders participating in smart-city music and art projects. This is achieved by embedding these protections within the vector-valued CyberRank system, which operates on a common basis of four shared dimensions: pay, voice, credit, and ecology . While the dimensional basis is universal, the framework's power lies in its ability to support role-specific kernel structures, allowing for tailored constraints and weights that reflect the unique contributions and vulnerabilities of each participant group . This nuanced approach avoids a rigid, one-size-fits-all solution, instead creating a flexible yet robust architecture for equitable value distribution and participation. The stakeholder roles to be considered include artists, musicians, curators, citizens, venue operators, and even AI systems acting as co-creators or facilitators .

The dimension of **Pay** is operationalized through the r_i^{pay} component of each stakeholder's fairness vector r_i . For artists and musicians, this translates into a guaranteed royalty share curve defined within their personal kernel K_i . This kernel specifies a hard lower bound on revenue, ensuring that no smart contract or commercial arrangement for a public XR performance can undercut their minimum earnings. The system would automatically reject any financial proposal that pushes their projected r_i^{pay} below this predefined floor, effectively codifying minimum wage or royalty agreements as a mathematical invariant . For venue operators and city entities, the pay dimension might be framed differently, perhaps focusing on budget adherence or return-on-investment thresholds, but it remains subject to the same viability logic. The key is that for every role, there is a defined minimum economic floor that the system is designed to protect.

The dimension of **Credit** is represented by the r_i^{credit} component and is tied to minimum visible attribution . In the world of collaborative and often algorithmically generated art, ensuring proper authorship is critical. The fairness kernel for an artist or developer would encode requirements for their name, logo, or unique identifier to be displayed alongside the work they contribute to, regardless of the final composition's success. This prevents exploitation where creators are rendered anonymous in commercially successful projects. The system could track attribution compliance through metadata embedded in the ALN, and any plan that results in a contributor being improperly uncredited would fail the viability check . For curators and platform operators, this dimension ensures that the provenance of all assets used in a curated experience is transparently documented, upholding intellectual property rights across the ecosystem.

Participatory **Voice**, captured by r_i^{voice}, is a more complex dimension to quantify but is central to the framework's goal of fostering genuine community engagement . For citizens participating in a public art project, r_i^{voice} could represent a guaranteed voting weight or governance rights over aspects of the experience, such as the selection of scenes, algorithmic curation rules, or the reuse of community-generated assets . This could be tokenized using a CybosToken-style internal rights system, where a citizen's contribution (e.g., participation hours, data sharing, or local knowledge) increases their voting power . Any governance proposal that dilutes creator or citizen voice below the level stipulated in their kernel would be automatically vetoed by the viability engine . For artists and musicians, their `r_i^{voice}` ensures a degree of co-ownership and veto power over how their work is used or modified in a collective piece, protecting against misattribution or inappropriate transformations. This dimension aims to shift power dynamics away from top-down control towards distributed, role-weighted governance.

Finally, the **Ecology** dimension, denoted `r_i^{ecology}`, binds contributors to measurable planetary benefits, making sustainability an intrinsic part of creative success in smart-city projects . This component links an artist's or team's success to ecological-sustainability metrics, such as local restoration scores, carbon offset contributions, or strict compliance with an energy budget allocated per event or XR node . For example, an XR installation powered by renewable energy sources might receive a higher ecological score than one drawing from a high-carbon grid. This incentivizes environmentally conscious creation and holds all participants accountable for the environmental footprint of their work. The system would measure these metrics through telemetric data from the XR grid and environmental sensors, feeding the results back into the fairness vector to determine if the project meets its ecological obligations . This dimension ensures that the pursuit of artistic expression does not come at an unacceptable environmental cost, creating a form of ecological co-ownership where the planet itself is considered a stakeholder .

By structuring the fairness model around this common four-dimensional basis but applying role-specific constraints, the framework achieves a sophisticated balance between universal principles and contextual pragmatism. It allows for a shared vocabulary of fairness while acknowledging that the stakes and responsibilities differ for an individual performer versus a city planner. Mathematically, this is treated as a family of related kernels, K^{role}, all residing within the same coordinate system, which preserves the ability to compare and reason about cross-role fairness and equity . This structure is the key to building a system that is not only provably safe and fair but also deeply integrated into the economic, social, and ecological realities of the communities it serves.

Integration with Existing Systems: The Vector CyberRank Upgrade Ring

A critical requirement for the successful implementation of this research framework is its seamless integration with the existing technological stack, specifically the legacy scalar

CyberRank system. The user has explicitly stated that the transition cannot be a hard cut-over; instead, the new vector CyberRank must function as an "upgrade ring" that remains subordinate to established controls and outputs . This backward compatibility is paramount to ensure that the introduction of this new, more sophisticated ranking system does not decouple any existing pipelines or temporarily weaken safety and rights guarantees . The recommended path is a phased, co-existence model where the vector rank acts as a powerful refinement signal, tightening decisions made by the primary scalar rank without ever overriding its core safety functions .

The initial implementation phase will involve generating both scalar and vector forms of CyberRank simultaneously . The legacy scalar CyberRank, which currently represents a probability of observation, will continue to serve its original purpose in routing and user experience (UX) . However, the new vector CyberRank, denoted as a vector $r \in \mathbb{R}^m$, introduces multiple axes of evaluation, including the newly added **stakeholder fairness** components alongside the existing ones like safety, legal robustness, biomech cost, psych risk, and ecological impact . This multi-axis evaluation fundamentally changes the nature of the decision-making process. Instead of choosing between two options based on a single score, the system can now evaluate trade-offs across many dimensions, preventing a situation where a project is deemed "good" by the scalar rank due to high entertainment value but is simultaneously unsafe, unfair, or ecologically damaging.

The decision logic for filtering candidate actions is structured as a layered gate, ensuring that no constraint is bypassed. The sequence begins with the most fundamental layer: **Safety Kernels**. All candidate plans are first checked against the controlled-invariant viability kernels to ensure they do not violate any hard-coded physical, mental-health, or safety boundaries . Plans that fail this initial viability check are immediately rejected . The second layer is the **Stakeholder Fairness Kernels**. Plans that pass the safety check are then evaluated against the role-specific fairness kernels for pay, voice, credit, and ecology . Again, any plan that would cause any stakeholder to fall below their guaranteed minimum floor is filtered out at this stage . Only after successfully passing through these two mandatory constraint layers do plans proceed to the third and final layer: **Capability Maximization**. Here, the full vector CyberRank profile is used to rank and select the optimal action. This is where Pareto algebra comes into play . Given a set of feasible plans that are all "Pareto-safe" (i.e., none widens any invariant set or weakens any right), the system selects the one that offers the best improvement across the vector of criteria, such as enhancing safety, fairness, and energy efficiency simultaneously . This Tsafe selection process ensures that among equally safe and fair options, the one with the highest overall "goodness" score is chosen .

This hierarchical filtering process makes the vector CyberRank a powerful governance tool. It does not replace the scalar rank but rather provides a deeper, more nuanced signal that informs a more responsible decision. The vector rank's influence can be gradual; initially, it may only be used to flag potentially problematic plans for human review. Over time, as the system's stability and predictive accuracy improve, the vector rank can be given more authority in the selection process, always operating under the explicit rule that it can never weaken existing rights or widen safety envelopes . This conservative, evolutionary approach to integration de-risks the deployment and builds confidence in the new system. It leverages the existing, battle-tested scalar rank as a stable anchor while incrementally adding the sophisticated capabilities of multi-axis evaluation. The ultimate goal is to create a hybrid system where the legacy UX driven by scalar CyberRank benefits from the hidden intelligence of the vector rank, leading the ecosystem toward patterns that are demonstrably safer, fairer, and more sustainable over time without requiring constant manual policing . The mathematical task of formalizing Pareto dominance and Tsafe selection within this constrained environment is a key research component, ensuring that the dynamics of the vector CyberRank process are well-posed and converge toward desirable outcomes .

System Layer	Primary Function	Governing Principle	Interaction with New Vector Rank
Legacy Scalar CyberRank	Routing, UX, and legacy observation probability .	Optimizes for a single heuristic score.	Continues to operate as the primary output for UX and routing. Remains the default if the vector rank is not yet fully trusted.
Controlled-Invariant Envelopes	Enforces hard lower bounds on all stakeholder dimensions .	Non-regression guarantee; no plan can push a stakeholder below their minimum floor.	Acts as a mandatory pre-filter. The vector rank has no effect if a plan fails this layer.
Vector CyberRank Controller	Multi-axis evaluation of plans based on safety, fairness, ecology, etc. .	Pareto-safe maximization; chooses the best option among those that respect all constraints.	Operates as a secondary filter and optimizer *after* the safety and fairness envelopes have been satisfied.
Decision Engine	Selects the final action from candidate plans.	Hierarchical filtering: Safety -> Fairness -> Capability Maximization .	Employs Tsafe selection to choose the optimal plan from the pool that passes all prior layers.

This carefully designed integration pathway ensures that the advanced mathematical structures of the fairness envelopes and vector CyberRank can be deployed immediately in a provably safe manner. It respects the existing infrastructure while paving the way for a future where intelligent systems make increasingly complex, ethically-aware decisions based on a comprehensive understanding of risk and value.

Practical Application in XR-Grid Environments

The theoretical constructs of controlled-invariant envelopes and vector CyberRank find their most immediate and impactful application in the dynamic, immersive environments of smart-city XR grids, particularly in contexts analogous to Phoenix, characterized by intense sunlight, high ambient temperatures, and potential resource constraints . In these settings, XR deployments are not merely digital experiences; they are physical interventions that interact with the urban environment, affecting everything from public safety and sensory comfort to energy consumption and ecological balance. The research framework provides a direct mechanism to manage these complex interactions, ensuring that creative projects are both innovative and responsibly grounded in their physical context. The immediate need is to bound tangible parameters like haptics, sound intensity, session length, crowd density, and ecological load at each node—from a small plaza to a major route—to make the creative experience provably safe, reversible, and sustainable .

For a Phoenix-like XR-grid, the environmental context heavily influences the constraints encoded within the fairness kernels. For example, the "energy use" dimension of the ecological kernel becomes critically important. High-energy-consuming XR experiences could strain the local power grid, especially during peak demand hours, and contribute to the urban heat island effect through increased waste heat from computing infrastructure. Therefore, the fairness kernel for a venue operator or a city entity would contain strict caps on kilowatt-hour consumption per hour of operation. The vector CyberRank would then reward projects that demonstrate superior energy efficiency, perhaps by leveraging edge AI to minimize cloud dependency or by using algorithms that adapt visual fidelity based on available solar power [10] [57]. Similarly, the "crowd density" parameter is vital for public safety. In a crowded plaza hosting a large-scale AR concert, the system would use telemetric data to monitor real-world pedestrian flow. If the XR deployment were to cause dangerous crowding, the viability kernel would trigger a corrective action, such as dynamically reducing the number of avatars or lowering the volume of spatial audio to encourage dispersal, thereby enforcing a maximum allowable density .

The sensory parameters of haptics and sound intensity are directly linked to the "mental-health impact" and "safety" components of the fairness vector for citizens . In a hot, arid climate, excessive haptic feedback could lead to skin irritation or overheating, while overly loud or disorienting soundscapes could cause stress or anxiety. The viability kernel for a citizen's sensory exposure would define hard limits on dB(A) levels and haptic force/timing profiles. These limits would be enforced in real-time by the Cybercore-Brain controller, which adjusts the XR experience based on ALN telemetry from participants' implants or devices. This ensures that an artist's ambitious sonic artwork cannot become a public nuisance or a health hazard. The system's ability to automatically correct for such violations embodies the principle of reversibility, allowing for creative exploration without irreversible negative consequences .

Furthermore, the framework supports novel forms of civic engagement and art. For instance, an ecological art project could be designed to tie directly into the `r_i^{ecology}` component of the fairness vector. An XR installation might visualize real-time air quality data from city sensors, with its aesthetic evolving based on pollution levels. Contributors to this project, including the artists and the citizens whose data feeds the system, would see their ecological credit score (`r_i^{ecology}`) improve as the project contributes to public awareness and potentially inspires behavioral changes that reduce pollution [2] [31] . The vector CyberRank would promote such projects over purely entertainment-focused ones that lack a positive ecological or social impact, subtly guiding the cultural landscape of the city towards more meaningful and sustainable ends. This aligns with the growing trend of using XR and AI to address urban challenges, from sustainable mobility to enhanced citizen engagement [3] [4] [52] . The framework provides the mathematical scaffolding to operationalize these goals, turning abstract aspirations for a "smarter" or "greener" city into concrete, enforceable rules governing the deployment of creative technology. By making these physical and ecological parameters first-class citizens in the decision-making process, the framework ensures that the evolution of the metaverse within our cities is guided by principles of care for both people and the planet [5] [58] .

Parameter	Associated Dimension(s)	Enforcement Mechanism	Example Application in Phoenix-like XR Grid
Sound Intensity	Pay, Credit, Ecology	Viability kernel sets max dB(A) limits. Real-time monitoring via ALN. Auto-correction (e.g., volume reduction) if limit is breached.	Prevents auditory overload in crowded plazas and reduces noise pollution impacting nearby residents and wildlife.
Haptic Feedback	Pay, Credit, Ecology	Viability kernel sets max force/duration limits. Corrective actions include dampening haptic pulses or reducing duration.	Avoids skin irritation or discomfort in high-heat environments; ensures accessibility for users with sensory sensitivities.
Crowd Density	Pay, Credit, Ecology	Telemetric tracking of real-world foot traffic. Kernel sets max occupancy. Actions include dynamic avatar scaling or spatial audio panning.	Manages public safety at popular XR venues, prevents dangerous bottlenecks, and ensures equitable access to space.
Session Length	Pay, Credit, Ecology	Kernel sets max continuous interaction time. System issues reminders or prompts for breaks.	Mitigates fatigue and potential psychological distress from prolonged immersion in intense XR experiences.
Energy Consumption	Pay, Credit, Ecology	Kernel sets max kwh/node/hour. Vector rank rewards projects using energy-efficient algorithms or renewable sources.	Reduces strain on the local power grid and mitigates the urban heat island effect from data center waste heat.

This practical application demonstrates that the framework is not an academic exercise but a direct response to the real-world challenges of deploying immersive technology in complex urban ecosystems. It provides a concrete methodology for balancing the immense creative potential of XR with the fundamental needs of public safety, environmental sustainability, and social equity.

Quantifying Abstract Guarantees: Measuring Voice and Ecological Footprint

A significant challenge in implementing the fairness framework is translating its abstract principles of "voice" and "ecological footprint" into concrete, measurable, and computable metrics. For the system to function as intended, these concepts cannot remain philosophical ideals; they must be expressed as explicit algebraic formulas or finite-sum rules suitable for inclusion in ALN telemetry and downstream analysis . The research must provide a clear methodology for defining these quantities, establishing their initial values, and continuously updating them based on system behavior and external evidence.

The dimension of **Voice** (`r_i^{voice}`) can be operationalized through a combination of participatory metrics and token-based governance. Conceptually, it represents a stakeholder's guaranteed influence in decision-making processes relevant to them. A computable proxy for this could be a weighted score derived from several factors. For a citizen participating in a neighborhood art project, their voice score could be a function of: 1. **Contribution Hours:** The total time spent in the XR environment contributing to the project (e.g., providing input, testing features). 2. **Proposal Submissions:** The number of governance proposals submitted regarding the project's content, curation, or rules. 3. **Vote Participation Rate:** The percentage of relevant votes (e.g., on new scene selections or algorithmic updates) in which they participate. 4. **Success Rate of Votes:** The historical success rate of their passed proposals, which could act as a multiplier to reward effective participation.

Mathematically, this could be represented as a weighted sum:

$$\text{VoiceScore}_i = w_1 \cdot H_i + w_2 \cdot P_i + w_3 \cdot V_i + w_4 \cdot S_i$$

where H_i is contribution hours, P_i is proposal count, V_i is vote participation rate, and S_i is the success rate of past votes for stakeholder i. The weights (w_j) could be adjusted by the system to prioritize certain types of engagement. This score would be stored in their fairness kernel and compared against a minimum threshold defined by the project's governance rules. This approach ties the abstract notion of "voice" to observable, recordable actions within the system. The reference to CybosToken-style rights suggests that this score could be linked to an internal cryptocurrency or reputation token, where a higher voice score grants greater voting power or influence in economic decisions, such as revenue sharing from a successful project .

The dimension of **Ecological Footprint** (`r_i^{ecology}`) requires a similarly rigorous definition to be useful. It must capture the net environmental impact of a stakeholder's activities within the XR grid. A multi-faceted formula for this metric could incorporate several key indicators: 1. **Energy Cost:** The total kilowatt-hours (kWh) of energy consumed by the stakeholder's hardware and the network resources supporting their project. This could be normalized per hour of activity or per unit of creative output. 2. **Carbon Equivalent:** A conversion of the energy cost into a CO2 equivalent, using regional grid emission factors. This provides a standardized measure of climate impact. 3. **Water Footprint:** For a Phoenix-like environment, this is particularly salient. It would account for the water used in cooling data centers or other physical infrastructure supporting the XR grid, measured in liters or cubic meters. 4. **Waste Generation:** A measure of e-waste generated over the lifecycle of the hardware used, or the digital waste from abandoned or poorly optimized XR assets. 5. **Positive Impact Credits:** Offsetting factors that could increase the ecological score, such as contributions to a verified reforestation project, use of 100% renewable energy, or the promotion of pro-environmental behaviors within the XR experience.

A simplified ecological score could be calculated as:

\text{EcologicalScore}<em _text_energy="\text{energy">i = C}} \cdot E_i + C_

where E_i, CO2_i, and W_i are the respective negative impacts for stakeholder i, and O_i is the positive impact credits they have earned. The coefficients (C) represent the relative weighting of each factor, which could be determined by city sustainability policies or project-specific goals. This score would be tracked in real-time via telemetry from the XR grid's infrastructure and environmental sensors. The fairness kernel would then establish a minimum acceptable ecological score; for example, a venue operator might be required to maintain a score above a certain threshold, with any shortfall triggering a corrective action, such as a temporary suspension of operations until improvements are made. This transforms the ecological dimension from a vague aspiration into a hard constraint that is actively managed and optimized by the system, ensuring that the creative economy is genuinely aligned with planetary boundaries [8] [10] [32] .}} \cdot W_i - C_{\text{offset}

By developing these explicit formulas, the research provides the necessary tools to make fairness a computable and enforceable property. The initial calibration of these formulas —for setting the weights and thresholds in the equations—would be a critical task, likely involving collaboration between data scientists, domain experts (urban planners, ecologists), and community representatives to ensure the metrics are both technically sound and socially meaningful. Once established, these metrics can be continuously refined and calibrated over time as the system gathers more longitudinal data, with the proviso that any updates can only tighten the envelopes and preserve existing rights, never weaken them .

Phased Implementation Roadmap and Future Research Directions

The successful deployment of the proposed fairness framework requires a pragmatic, phased implementation roadmap that prioritizes immediate, high-impact tasks while laying the groundwork for more complex, long-term research. The user has provided a clear strategic direction, separating the immediate priorities—controlled-invariant envelopes and vector-valued CyberRank—from the longer-term theoretical work on jurisdiction lattices and F–H–E functionals . This approach de-risks the project by delivering tangible safety and fairness guarantees in the short term, while keeping the door open for future enhancements that build upon this solid foundation.

The immediate, highest-priority action is the development and integration of **controlled-invariant fairness envelopes**. This involves specifying new ALN particles, such as `fairness.kernel.stakeholder.v1`, which will carry the definitions of each stakeholder's multi-dimensional envelope . Concurrently, Rust crates must be developed to encode these kernels as polytopes and to compute the critical "kernel distance" and "fairness-exit risk" metrics in real-time from telemetry data . This work is the bedrock of the entire system; without provable safety and ecological bounds, no other fairness mechanisms can be reliably deployed. The outputs of this phase will be directly applicable to bounding physical and ecological parameters in smart-city XR projects, such as haptic intensity, sound pressure levels, and energy consumption, making creative experiences provably safe and reversible from day one .

Following the establishment of the envelope geometry, the second priority is the implementation of the **vector-valued CyberRank system**. This involves creating a new particle, `cyberank.vector.profile.v2`, to hold the multi-axis rankings . The key challenge here is designing the backward-compatible integration path with the legacy scalar CyberRank. The system must be designed to generate both scores initially, with the vector rank acting as a refinement signal that tightens decisions made by the scalar rank, never weakening them . The decision logic must be formalized to filter candidate actions through a hierarchy of safety kernels, stakeholder fairness kernels, and finally, capability maximization using Pareto algebra . This phase will enable the system to begin ordering and promoting projects based on a holistic view of safety, fairness, and ecological impact, rewarding "good" patterns by default without stifling creativity .

Once these two core components are functioning and providing tangible benefits in smart-city projects, the research can progress to the more abstract, long-term goals that were explicitly deferred. The first of these is **Jurisdiction and Policy Lattice Theory**. With a working model of stakeholder fairness, the next step is to formalize how these fairness kernels interact with the patchwork of legal and regulatory jurisdictions across different regions . This involves modeling policy profiles as elements of a complete lattice, where the lattice operations (meet and join) correspond to harmonization rules like "strictest-wins" . This would yield a globally reusable way to reason about compatible vs. incompatible rules, helping to resolve conflicts between different governance stacks and bound legal risk for system trajectories .

The final deferred area is the development of **Free-energy / harm trade-off functionals**. This research aims to make the estimators for free-energy drop (F), ethical/somatic harm (H), and dissipation/energy cost (E) fully explicit and computable from particle fields and telemetry . This would provide a universal thermodynamic-like framework for comparing the efficiency and safety of different intelligent systems and actions . Proving that controllers minimizing a composite F-H-E objective cannot cross soul.guardrail envelopes would create a powerful link between information-theoretic efficiency and ethical safety . This work, while foundational, depends on the empirical data and validated models generated during the earlier phases of implementation.

This phased roadmap ensures that the research remains focused and deliverable. It starts with solving the most pressing, practical problems of safety and fairness in public XR deployments. As the system matures and generates more data, it can tackle the more complex challenges of global governance and fundamental physics-based safety constraints. Throughout this journey, the overarching principle of spreadsheet-grade, compositional mathematics will guide the development, ensuring that every component remains transparent, analyzable, and ultimately, deployable within the Cybercore-Brain ecosystem .

Reference

1. State-of-the-Art on eXtended Reality Streaming https://hal.science/hal-04977125/file/HAL___XR_streaming_survey.pdf
2. Metaverse in Smart Cities: Transforming Urban Life and ... https://www.researchgate.net/publication/398076115_Metaverse_in_Smart_Cities_Transforming_Urban_Life_and_Governance
3. Artificial Intelligence in Urban Planning: A Bibliometric ... https://www.mdpi.com/2073-445X/14/11/2100
4. A Data-Driven Systematic Review of the Metaverse in ... https://www.sciencedirect.com/science/article/pii/S152614922500270X
5. Intelligent World 2030 https://www.huawei.com/admin/asset/v1/pro/view/d2c1c28eeba24f4ca7bdf0022805a1dc.pdf
6. Designing Interactions with Generative AI for Art and ... https://dl.acm.org/doi/10.1145/3715336.3735843
7. CVPR 2025 Accepted Papers https://cvpr.thecvf.com/Conferences/2025/AcceptedPapers
8. A Review from Technical Landscape to Regulatory Specific https://arxiv.org/pdf/2503.03539
9. Hybrid Rule-Based and Reinforcement Learning for Urban ... https://www.mdpi.com/2076-3417/15/19/10761
10. Green IoT and Edge AI as Key Technological Enablers for a ... https://pmc.ncbi.nlm.nih.gov/articles/PMC8434294/
11. THE VIRTUAL PRODUCTION FIELD GUIDE https://cdn2.unrealengine.com/vp-field-guide-v1-3-01-f0bce45b6319.pdf
12. glove.6B.100d.txt-vocab.txt https://worksheets.codalab.org/rest/bundles/0xadf98bb30a99476ab56ebff3e462d4fa/contents/blob/glove.6B.100d.txt-vocab.txt
13. download the full English word list here https://www.freescrabbledictionary.com/english-word-list/download/english.txt
14. Computer Science Curricula 2013 https://www.acm.org/binaries/content/assets/education/cs2013_web_final.pdf
15. Music, song and the performing arts tradition: past ... https://unesdoc.unesco.org/ark:/48223/pf0000134331

16. Immersive Multimedia Communication: State-of-the-Art on ... https://dl.acm.org/doi/10.1145/3721292
17. Arxiv今日论文| 2025-12-11 - 闲记算法 http://lonepatient.top/2025/12/11/arxiv_papers_2025-12-11
18. A Comprehensive Survey of Machine Learning Applied to ... https://ieeexplore.ieee.org/iel8/9739/11321210/10930485.pdf
19. Artificial Intelligence 2025 https://www.arxiv.org/list/cs.AI/2025?skip=28975&show=2000
20. Toward Intelligent AIoT: A Comprehensive Survey on ... https://www.mdpi.com/2227-7390/13/21/3382
21. 人工智能2025_7_16 http://arxivdaily.com/thread/69481
22. 首届中英三校博士生联合论坛 暨2025年 ... https://www.xjtlu.edu.cn/wp-content/uploads/2025/04/2025-Symposium-Abstract-Brochure.pdf
23. An Experiment in Method. In 110th CAA Annual Conference https://www.academia.edu/73674173/Phantom_Images_Residual_Violences_An_Experiment_in_Method_In_110th_CAA_Annual_Conference
24. Computer Science Jun 2025 https://www.arxiv.org/list/cs/2025-06?skip=8975&show=1000
25. 斯坦福CS课程汇总_karupinwm. - 华为开发者空间 https://huaweicloud.csdn.net/64ddeaee4c7ead5211f11b85.html
26. 333333 23135851162 the 13151942776 of 12997637966 https://www.cs.princeton.edu/courses/archive/spring25/cos226/assignments/autocomplete/files/words-333333.txt
27. Allwords Withdigits | PDF | Breastfeeding https://www.scribd.com/document/574115973/allwords-withdigits
28. 2024 - Sustainability Report - SF Express https://ir.sf-express.com/media/pljplgmn/%E9%A1%BA%E4%B8%B0%E6%8E%A7%E8%82%A12024%E5%B9%B4%E5%BA%A6%E5%8F%AF%E6%8C%81%E7%BB%AD%E5%8F%91%E5%B1%95%E6%8A%A5%E5%91%8A-%E8%8B%B1%E6%96%87.pdf
29. China Construction Bank Corporation 2024 Sustainability ... https://www.ccb.com/cn/html1/office/ebank/dsh/CCB%20Sustainability%20Report%202024.pdf
30. China Overseas Land & Investment Limited https://www1.hkexnews.hk/listedco/listconews/sehk/2025/0428/2025042801201.pdf
31. Strengthening and Implementing the Global Response https://www.ipcc.ch/site/assets/uploads/sites/2/2022/06/SR15_Chapter_4_HR.pdf
32. China https://www.nsfc.gov.cn/Portals/0/fj/fj20230220_01.xlsx

33. Amps Proceedings Series 32 https://amps-research.com/wp-content/uploads/2024/02/Amps-Proceedings-Series_32.pdf
34. Human-Centered Design, Operation and Evaluation of Mobile ... https://link.springer.com/content/pdf/10.1007/978-3-031-93064-5.pdf
35. International Symposium on Machine Learning and Art 2021 https://www.cityu.edu.hk/scm/artmachines2/AM2%20Conference%20Proceedings.pdf
36. Jaes V52 6 All PDF https://www.scribd.com/document/481617939/JAES-V52-6-ALL-pdf
37. Multi Text8 E10 d300 Vs2e-4 Lr1e-5 Margin1.words | PDF https://www.scribd.com/document/480950803/multi-text8-e10-d300-vs2e-4-lr1e-5-margin1-words-txt
38. Computer Science https://www.arxiv.org/list/cs/new?skip=200&show=2000
39. Design Tools and Methods in Industrial Engineering IV https://link.springer.com/content/pdf/10.1007/978-3-031-76594-0.pdf
40. r2023 Bme Bbe | PDF | Engineering | Biomedical ... https://www.scribd.com/document/971688860/r2023-Bme-Bbe
41. Computer Science Curricula 2023 - CS2023 - ACM https://csed.acm.org/wp-content/uploads/2023/09/Version-Gamma.pdf
42. Proceedings | CHI 2021 - ACM https://chi2021.acm.org/proceedings
43. Body of Knowledge - CS2023 https://csed.acm.org/wp-content/uploads/2024/01/Body-of-Knowledge-v1-bookmarksv2.pdf
44. Blog https://interactions.acm.org/blog/view/addressinginstitutional-racism-within-initiatives-forsigchis-diversity-an
45. Encyclopedia of Remote Sensing (Encyclopedia of Earth ... https://www.scribd.com/document/341391859/Encyclopedia-of-Remote-Sensing-Encyclopedia-of-Earth-Sciences-Series-2014
46. Running a museum: a practical handbook https://unesdoc.unesco.org/ark:/48223/pf0000141067
47. sd2023.pdf https://www.chinamobileltd.com/en/ir/reports/ar2023/sd2023.pdf
48. Conference proceedings https://unesdoc.unesco.org/ark:/48223/pf0000149641
49. Computer Science Curricula 2013 - Algorithms, 4th Edition https://algs4.cs.princeton.edu/home/CS2013-final-report.pdf
50. Mouser Methods V4i3 | PDF | Augmented Reality https://www.scribd.com/document/710008702/Mouser-Methods-v4i3
51. Localization as a Key Enabler of 6G Wireless Systems https://ieeexplore.ieee.org/iel7/8782661/10008219/10287134.pdf
52. Artificial Intelligence for the Metaverse: A Survey https://arxiv.org/pdf/2202.10336

53. AI-enabled routing in next generation networks: A survey https://www.sciencedirect.com/science/article/pii/S111001682500122X
54. Control Charts and Machine Learning for Anomaly ... https://hal.science/hal-03561200/file/book%20TRAN.pdf
55. Reconfigurable Intelligent Surfaces (RIS) and Their Role in ... https://onlinelibrary.wiley.com/doi/10.1002/ett.70164
56. 2024 Argentine Conference on Electronics https://ieeexplore.ieee.org/iel7/10487127/10487112/10487120.pdf
57. Word https://www.itu.int/dms_pub/itu-t/md/22/tsag/td/240122/GEN/T22-TSAG-240122-TD-GEN-0409!!MSW-C.docx
58. Recent Trends of AI Technologies and Virtual Reality https://link.springer.com/content/pdf/10.1007/978-981-96-1154-6.pdf
59. Sohungomit Urban Drakazan | PDF https://www.scribd.com/document/761912805/Sohungomit-Urban-Drakazan
60. (PDF) Proceedings of the 21st International Conference on ... https://www.academia.edu/69113002/Proceedings_of_the_21st_International_Conference_on_Auditory_Display_ICAD_2015_Graz_Austria_July_8_10_2015
61. words_SG_upto2020.txt https://zenodo.org/record/5516252/files/words_SG_upto2020.txt
62. Simulation, Modeling, and Programming for Autonomous ... https://link.springer.com/content/pdf/10.1007/978-3-319-11900-7.pdf
63. MetaWatch https://ieeexplore.ieee.org/iel8/6488907/6702522/10998456.pdf
64. (PDF) Split Federated Learning for 6G Enabled-Networks https://ieeexplore.ieee.org/iel7/6287639/6514899/10385040.pdf
65. Blog https://interactions.acm.org/blog/view/the-ux-ownership-war
66. Technical program for Friday June 6, 2014 https://ieeexplore.ieee.org/iel7/6849600/6858556/06858575.pdf
67. Computer Science Curricula 2023 https://csed.acm.org/wp-content/uploads/2023/03/Version-Beta-v2.pdf
68. Body of Knowledge - CS2023 https://csed.acm.org/wp-content/uploads/2024/01/Body-of-Knowledge-v1-1.pdf
69. (PDF) URLLC for 5G and Beyond: Requirements, Enabling ... https://ieeexplore.ieee.org/iel7/6287639/6514899/09406015.pdf
70. Student Poster Book of Abstracts https://ewh.ieee.org/soc/pes/sasc/file/Book_of_Abstracts_Student_Poster_GM14.pdf

71. Draft management plan for Petra archaeological and ... https://unesdoc.unesco.org/ark:/48223/pf0000094949
72. List of eligible projects - European Innovation Council https://eic.ec.europa.eu/document/download/c0a62cd5-8262-4b5b-865a-c2a007491fc1_en?filename=List%20eligible%20Projects%20EIC%20Transition%202024.xlsx
73. Computer Science https://www.arxiv.org/list/cs/new?skip=25&show=1000
74. Conference Proceedings https://www.researchgate.net/profile/Hameeza-Ahmed-2/publication/356790323_Towards_Efficient_Vectorization_using_Polyhedral_Compilation/links/61adb90cca2d401f27cb90d5/Towards-Efficient-Vectorization-using-Polyhedral-Compilation.pdf
75. China https://www.nsfc.gov.cn/Portals/0/fj/fj20240205_01.xlsx
76. Springer Series on Cultural Computing Founding Editor https://www.academia.edu/123963596/Springer_Series_on_Cultural_Computing_Founding_Editor
77. Blog https://interactions.acm.org/blog/view/why-codesigning-ai-is-differentand-difficult
78. Artificial Intelligence in HCI https://link.springer.com/content/pdf/10.1007/978-3-031-93429-2.pdf
79. Arxiv今日论文| 2025-11-20 - 闲记算法 http://lonepatient.top/2025/11/20/arxiv_papers_2025-11-20
80. Artificial Intelligence https://arxiv.org/list/cs.AI/new
81. Wednesday, March 23 Thursday, March 24 https://ieeexplore.ieee.org/iel7/9764877/9764959/09765086.pdf
82. Information, Volume 16, Issue 7 (July 2025) – 110 articles https://www.mdpi.com/2078-2489/16/7
83. from algorithms to socio-technical analyses on AI https://theses.hal.science/tel-04963643v1/file/142478_DELARUE_2025_archivage.pdf
84. Program Book 10-18 | PDF | Mathematical Optimization https://www.scribd.com/document/785837327/Program-Book-10-18
85. 5 https://worksheets.codalab.org/rest/bundles/0xd74f36104e7244e8ad99022123e78884/contents/blob/frequent-classes
86. Treasure trail guide osrs 2 factor guide https://movefamunumuxe.weebly.com/uploads/1/4/1/9/141968622/zoberetimifid.pdf
87. WIRELESS PERSONAL COMMUNICATIONS - Springer Link https://link.springer.com/content/pdf/10.1007/978-1-4615-6237-5.pdf
88. Rats, mice and people: https://www.researchgate.net/profile/Rhodes_Makundi/publication/260025210_Interaction_between_rodent_species_in_agro-

forestry_habitats_in_the_western_Usambara_mountains_northeast_Tanzania_and_its_potential_for_plague_transmission_to_humans/links/0046352f2125b7eeee000000.pdf

89. Download book PDF - Springer Link https://link.springer.com/content/pdf/10.1007/978-4-431-68325-4.pdf
90. How to Open a Grassroots Music Venue https://musicvenuetrust.com/wp-content/uploads/2019/07/How-to-Open-a-Grassroots-venue.pdf
91. Amendment No. 4 to Form F-1 https://www.sec.gov/Archives/edgar/data/1610601/000119312514438736/d816932df1a.htm
92. hw3_stats_google_1gram.txt https://www.cs.cmu.edu/~roni/11661/2017_fall_assignments/hw3_stats_google_1gram.txt
93. Experts for Sk Slovenia Huntington Association https://www.linknovate.com/search/?query=sk%20slovenia%20huntington%20association
94. Steering Histories: A Scientifically Rigorous Framework for ... https://www.researchgate.net/publication/394883822_Steering_Histories_A_Scientifically_Rigorous_Framework_for_Temporal-Navigation_Research_Devices