Quantifying Creative Integrity and Social Impact: A Cybernetic Framework for Measuring Fairness in AI-Driven Entertainment

This report provides a comprehensive analysis of a proposed research agenda aimed at developing and empirically validating a fairness-centered AI scoring framework for the entertainment industry, with an initial focus on music. The core innovation of this framework is the treatment of abstract ethical constructs—such as integrity, harm, and impact—as concrete, measurable entities encoded within Cybernet Knowledge Objects (KOs). This approach seeks to move beyond theoretical discourse into a domain of empirical validation, technical implementation, and transparent governance. The research prioritizes the development of foundational KOs related to integrity-weighted tool use and harm-hype trade-offs, while concurrently addressing critical risks of systemic bias and flawed causal attribution. The ultimate objective is to create a system where AI assistance and content routing are guided by evidence-based metrics that promote long-term creator support, listener trust, and equitable outcomes across diverse communities.

Foundational Cybernet Knowledge Objects: From Abstract Principles to Technical Specifications

The cornerstone of the proposed framework is the transformation of abstract ethical goals into tangible, implementable artifacts known as Cybernet Knowledge Objects (KOs). These KOs serve as structured data containers that encode specific, measurable properties of creative works, enabling them to be processed by AI systems for purposes of scoring, routing, and assistance. The initial phase of this research focuses on three pivotal KOs: `KO.IntegrityTrace.music.v1`, `KO.ImpactLog.music.v1`, and `KO.FairRoutingExperiment.v1`. Each object is meticulously designed to capture distinct dimensions of a work's lifecycle, from creation to societal reception, providing the raw material for calculating the core scores of integrity (`I`) and social impact (`S`). The design philosophy emphasizes starting with the music domain to build stable, validated metrics before generalizing the schema to other media like film and gaming, ensuring cross-domain comparability [22]. This initial focus on music tracks, music videos, and audio deployed in public spaces in Phoenix allows for controlled, real-world experimentation before broader rollout [132].

The first foundational KO, `KO.IntegrityTrace.music.v1`, is engineered to quantify the integrity of a creative work. Its primary function is to calculate an `Integrity_score` ($I \in [-1,1]$) derived from three core signals: `AI_use_disclosed`, `plagiarism_risk`, and `marketing_truthfulness`. The signal `AI_use_disclosed` measures whether the utilization of artificial intelligence in the work's creation was explicitly stated by the artist or their representatives. This directly responds to strong consumer demand for transparency; a survey by the International Federation Of The Phonographic Industry (IFPI) revealed that 73% of music consumers agree that an AI system should clearly list any music it has ingested or used for training [2]. Furthermore, 76% believe an artist's music or vocals should not be used by AI without permission, indicating a clear market expectation for accountability [2]. The recent tightening of rules by platforms like Spotify on AI voice impersonation further validates this trend towards greater disclosure [145189]. The second signal, `plagiarism_risk`, assesses the likelihood that a piece contains plagiarized melodic or structural elements. This is informed by emerging technologies in Music Information Retrieval (MIR), such as MelodySim, an audio-based similarity model designed to detect melody-related plagiarism by preserving the main melody while altering instrumentation and tempo [1] [86]. Another advanced method, TruMuzic, employs a deep learning system that analyzes musical notes, pitch, census metadata, and provenance data (the history and source information about a song and its creators) to achieve high accuracy in detecting various forms of plagiarism [26]. The third signal, `marketing_truthfulness`, evaluates the alignment between promotional claims made prior to release and the final artistic product. This presents a significant analytical challenge, likely requiring a multi-modal approach combining Natural Language Processing (NLP) to analyze press releases and social media posts against the final audio/video content, potentially augmented by human review for high-stakes cases. The relative weighting of these three signals within the final `Integrity_score` is not predetermined but will be a critical parameter to be determined through empirical validation via A/B testing [30].

The second KO, `KO.ImpactLog.music.v1`, is designed to measure the net social impact of a work, formalized as $S = B - H$, where `B` represents benefits and `H` represents harms. This construct aligns with established principles of risk management in AI, such as those outlined in standards like ISO/IEC 23894:2023 [65]. The `ImpactLog` proposes to capture incidents and their associated `H` (harm) and `B` (benefit) values, normalized by exposure counts to provide a rate per listener/viewer. A key feature of this KO is the inclusion of age bands, acknowledging that the impact of content varies significantly across different developmental stages. This is strongly supported by research linking adolescent mental health to digital content, including the potential for algorithms to push extreme content to vulnerable youth [40], and the unique cognitive and psychological vulnerabilities of children's brains to media influence [39] [127] [154]. The log would record incidents categorized by age group, allowing for more nuanced routing policies—for instance, stricter thresholds for public youth spaces versus more permissive but still evidence-based thresholds for adult-only contexts [48]. Quantifying benefits (B) is a particularly innovative and challenging aspect of this KO. Traditional recommender systems often implicitly assume B is near zero unless there is a clear commercial success metric like streams or sales [99]. The proposal to actively capture underreported positive outcomes is crucial for a balanced assessment. Research has demonstrated numerous beneficial effects of music engagement, including reduced symptoms of depression and anxiety, stress management, enhanced mood, and improved cognitive skills [11] [13] [14] [15]. To quantify B, the framework suggests using telemetry, opt-in studies, and sentiment analysis of user-generated content to estimate the probability of positive outcomes (P), leading to a benefit index of $B = \min(1, \beta P)$ [117]. This probabilistic approach draws inspiration from risk assessment modeling techniques [129].

The third KO, `KO.FairRoutingExperiment.v1`, serves as the experimental control layer for the entire framework. It is not a passive log but an active instrument for testing hypotheses about how different scoring configurations affect system behavior. Its purpose is to enable A/B tests where the intensity of AI assistance and the ranking of content are driven by different weightings of the core scores: creativity (C), integrity (I), and social impact (S). The `helpfulness_weight` ($H_w = w_C C + w_I I + w_S S$) becomes a decision surface that dictates the level of extra help an AI provides to a creator [30]. By systematically varying the weights w_C, w_I, and w_S, researchers can empirically measure downstream outcomes such as safety incident rates, creator diversity, and satisfaction among both small and large artists [30]. This KO also governs routing policies, such as using a reputation score ($R_{work} = W_t(w_B B - w_H H)$) to down-rank harmful works while preserving the visibility of creatively valuable but potentially controversial pieces [30]. The `FairRoutingExperiment` KO is essential for validating the entire framework; without it, the scores would remain untested hypotheses. It provides the mechanism to check for unintended consequences, such as the suppression of benign but low-impact genres when harm is heavily weighted [30]. The table below outlines the core components of these foundational KOs.

KO Name	Primary Purpose	Key Fields / Signals	Domain Specificity
`KO.IntegrityTrace.music.v1`	Quantify the honesty and originality of a creative work.	`AI_use_disclosed` (Boolean), `plagiarism_risk` (Float, [0,1]), `marketing_truthfulness` (Score, [-1,1])	Music-first, with schema designed for generalization to film/ gaming.
`KO.ImpactLog.music.v1`	Log age-banded harm and benefit incidents with exposure data.	`exposure_count` (Integer), `age_band` (String), `incident_type` (Enum), `severity` (Float), `event_timestamp` (DateTime)	Music-first, with schema designed for generalization to film/ gaming.
`KO.FairRoutingExperiment.v1`	Enable A/B testing of routing policies and AI assistance intensity.	`experiment_id` (String), `policy_config` (JSON), `variant_group` (String), `metrics_log` (JSON), `CHAT_spend` (Float)	Platform-agnostic, focused on policy logic rather than content specifics.

These KOs collectively form a cybernetic loop: the `IntegrityTrace` and `ImpactLog` provide the sensory data on a work's properties, the `repscore` synthesizes this data into a single reputation value, and the `FairRoutingExperiment` uses this value to generate commands (routing decisions, assistance levels) that are then fed back into the ecosystem, influencing future states and creating new data for the logs. This closed-loop system is the engine of the proposed fairness framework.

Empirical Validation of Integrity and Harm Metrics

The transition of the conceptual KOs into a functional framework hinges on rigorous empirical validation. The primary research thrust involves moving beyond theoretical definitions of integrity and impact to develop and test concrete measurement methodologies grounded in data. This process requires a multi-pronged approach combining signal development, quantitative modeling, and systematic A/B testing of the resulting policies. The initial empirical runs will be constrained to the specified domains —music tracks, music videos, and audio in XR/public-space deployments in Phoenix—to ensure data consistency and allow for targeted intervention and analysis [132].

The validation of the `KO.IntegrityTrace.music.v1` begins with the individual signals. For `plagiarism_risk`, the research can leverage existing MIR technologies. Models like MelodySim, which uses a Triplet Neural Network with a MERT encoder to compare melodic embeddings, offer a direct path to implementation [1]. The model's architecture, which computes a similarity matrix between all 10-second segments of two songs and flags them if over 20% of segments are classified as plagiarized, provides a replicable methodology [1]. The validation process would involve training this model on a curated dataset of original, sampled, and plagiarized tracks, including the novel MelodySim dataset with its carefully constructed variations that preserve melody while altering other musical aspects [86]. Similarly, the TruMuzic approach, which integrates multiple data types including musical notes, pitch, census data, and provenance data, offers a more holistic model whose features can be adapted [26]. Provenance data, which includes a trustworthiness score for the composer based on their history of collaborations and copying patterns, is particularly relevant for building a longitudinal integrity profile for both tracks and artists [26]. The validation of the `AI_use_disclosed` signal is more straightforward, relying on structured data from artist profiles, credits, and platform disclosures. The most complex signal to validate is `marketing_truthfulness`. This will require a combination of automated and manual processes. NLP models can be trained to analyze the semantic distance between promotional text (from press releases, social media, interviews) and lyrical themes or sonic characteristics of the final work. However, due to the subjective nature of many marketing claims, this will necessitate a tiered approach, with automated scoring for clear-cut cases and a workflow for human reviewers to adjudicate ambiguous ones. The final `Integrity_score` will be a weighted sum of these three signals, and the optimal weights will be determined through the experiments facilitated by `KO.FairRoutingExperiment.v1`.

The empirical validation of the `KO.ImpactLog.music.v1` centers on modeling the social impact equation $S = B - H$. Validating the harm component (H) requires establishing reliable incident reporting mechanisms. This goes beyond simple flagging to include age-banded logging. The framework must define what constitutes an "incident" (e.g., harassment reports, promotion of self-harm, incitement to violence) and link it to specific works via exposure data. The challenge lies in collecting sufficient exposure data to normalize incident rates, a problem noted in epidemiological studies of internet use [21]. The use of anonymized, consented telemetry from listening devices and platforms will be critical. For benefits (B), the validation is even more nascent. The formula $B = \min(1, \beta P)$, where P is the probability of a positive outcome, requires defining and measuring P [117]. Positive outcomes could range from documented prosocial behavior changes following music exposure interventions [10] to improvements in mood and cognitive function observed in older adults [14] and reductions in stress [15]. Capturing these underreported benefits may involve deploying opt-in studies where users consent to sharing biometric data (e.g., heart rate variability, EEG signals that can measure neural oscillations underlying emotion [16] [18]) or detailed diary entries during and after listening sessions. The coefficient β would then be calibrated to reflect the relative importance of different benefit types. The entire impact model must be tested against real-world outcomes. Does a higher calculated `Impact_score` correlate with increased long-term listener trust and sustained support for the creator? The `FairRoutingExperiment.v1`

KO is the essential tool for answering this question, allowing for controlled tests of routing policies that prioritize different combinations of creativity, integrity, and impact [191].

A central part of this validation effort is the concept of "harm weighs more than hype." This principle must be translated into a mathematical constraint. The research aims to fit functions for social impact and show parameter regions where short-term commercial success (e.g., click spikes, viral trends) would incorrectly outweigh evidence of long-term harm if not properly capped [30]. This involves stress-testing the reputation score formula $R_{work} = W_t(w_B B - w_H H)$ with historical data on trending but problematic content. For example, datasets showing the spread of disinformation networks influenced by recommendation algorithms could inform the setting of the harm weighting (w_H) [69]. The time-decay factor $W_t = e^{-\lambda T}$, which models the decay of a work's reputation over time, must also be empirically tuned [30]. Tracking the career trajectories of creators who have experienced public incidents, alongside anonymized telemetry on their productivity and well-being, can help identify an appropriate decay constant (λ) [30]. This ensures that severe but rare incidents have a longer memory, while lesser harms fade once the creator demonstrates behavioral improvement. This dynamic reputation system is designed to balance punitive measures with opportunities for forgiveness and redemption, a critical aspect of fair governance [178].

Mitigating Unintended Consequences: Bias Audits and Attribution Robustness

While the primary goal of the framework is to enhance fairness by penalizing harmful content, a sophisticated governance system must anticipate and mitigate the secondary harms that can arise from its implementation. The concurrent prioritization of Topic 4 (Bias & Inequality Audits) and Topic 7 (Incident Attribution Robustness) reflects a mature understanding of this challenge. Without these guardrails, the pursuit of a safer platform could inadvertently lead to the suppression of minority voices, the reinforcement of existing cultural biases, or the unjust stigmatization of creators based on flawed causal assumptions. These topics are not merely supplementary; they are essential for building a framework that is robust, defensible, and truly equitable.

Topic 4, Bias & Inequality Audits, directly confronts the risk that an `Impact_score`-based reputation system could unfairly disadvantage certain groups.

Recommender systems are known to suffer from popularity bias, where less popular items and creators are systematically disadvantaged, a phenomenon that can be exacerbated when safety metrics are introduced [148171172]. An algorithm might disproportionately associate lower-income or minority communities with higher incident rates due to biased policing and reporting practices, leading to a feedback loop where content from these communities is unfairly down-ranked [33]. To counteract this, the research proposes a systematic auditing process. This involves comparing incident rates *per unit of exposure* across different demographic, regional, and genre-based cohorts [30]. If the data shows that a particular genre or region has a higher raw incident count but also a much larger audience, the system can apply bias-correction weights to ensure the normalized rate reflects true causal impact rather than an artifact of enforcement disparities [30]. A key innovation proposed is the Minority Genre Protection Index (MGPI). This metric would serve as a real-time monitor, checking whether low-mainstream but culturally important genres are being under-promoted at specific `repscore` thresholds [30]. The MGPI would force the system to explicitly evaluate its own outputs for inequitable suppression and trigger compensation terms to maintain safety without erasing cultural diversity [30]. This approach draws parallels to historical solutions for technological displacement, such as the Record-Royalty Fund established in the 1940s to support live musicians displaced by recording technology, suggesting that a collective fund could be a viable model for supporting creators negatively impacted by the AI system [43].

Topic 7, Incident Attribution Robustness, addresses one of the most profound challenges in digital governance: assigning responsibility. Attributing complex societal outcomes, such as a rise in anxiety among a demographic or a coordinated protest, to a single piece of music or entertainment is fraught with peril. It opens the door to moral panics, censorship, and the unfair targeting of creators [91]. The research must therefore develop rigorous causal inference tools to distinguish mere correlation from actual causation [97]. The proposed solution involves building models that incorporate multiple lines of evidence: granular exposure data to confirm that the target population consumed the work, analysis of prior behavior to establish a baseline, and the use of control groups to isolate the effect of the work itself [30]. This statistical rigor is akin to the principles of algorithm assurance, which supports risk management for high-stakes algorithmic applications [98]. The output of these models should not be a definitive verdict but a probabilistic attribution with a confidence interval, quantifying the uncertainty inherent in the analysis [30]. Crucially, the system must define minimum evidential thresholds before an incident can be formally logged and used to alter a work's `Impact_score` or `repscore`. This involves quantifying the acceptable rates of false positives (wrongly blaming a work) and false negatives (failing to attribute a genuine harm) and designing the system to operate safely within these bounds [30]. This commitment to attributional humility is essential for preventing the framework from becoming a tool of ideological persecution or suppressing legitimate artistic expression, such as protest messages that may be flagged as "harmful" by overly simplistic classifiers [22]. The development of these causal tools is non-negotiable; without them, the `ImpactLog` would be built on a foundation of anecdotal and potentially erroneous data, undermining the entire credibility of the fairness framework.

Governance and Strategic Implementation Across Domains

The successful deployment of this AI scoring framework extends beyond technical implementation to encompass a robust governance structure and a clear strategy for expansion. The research plan wisely distinguishes between internal technical deliverables, intended for engineers and developers, and external policy-facing documents, designed for regulators, city partners like those in Phoenix, and the public [132]. This dual-layered approach ensures that the system is both practically buildable and transparently accountable. The framework's design principles are already aligned with major regulatory trends, particularly the European Union's Artificial Intelligence Act (AI Act), which establishes a precedent for governing high-risk AI systems [101][102].

Internally, the primary deliverable will be a set of detailed technical specifications. This includes precise JSON schemas for each KO, defining every field, its data type, and its constraints. It will contain the exact mathematical formulas for calculating the `Integrity_score` and `Impact_score`, along with the parameters for the A/B testing configurations in `KO.FairRoutingExperiment.v1`. Clear logging requirements for incident data, exposure metrics, and experiment outcomes will be mandated to ensure data quality and reproducibility. These specifications are the blueprint for integrating the framework into the Cybernet and Reality.os ecosystems, enabling the creation of a self-regulating system where AI assistance and content routing are dynamically adjusted based on empirically validated scores [77][122]. The ultimate integration point for these scores is the `knowledge_factor`, which determines the reliability and utility of KOs within the system. Exploring bounded couplings where the knowledge factor is weakly influenced by the social impact score ($F_K^{ent} = F_K + \kappa S_i$) represents a deep cybernetic principle: a system's overall intelligence and trustworthiness could be elevated by the quality of the content it processes, creating a virtuous cycle where positive and honest contributions are rewarded with greater systemic influence [103].

Externally, the research will produce governance frameworks and impact simulations. These documents translate the technical specifications into a narrative about the system's philosophy, its ethical trade-offs, and its expected societal effects. Impact simulations are a powerful tool for engaging stakeholders who are not fluent in code [75]. By visualizing the outcomes of different policy choices—such as adjusting the weights w_C, w_I, and w_S—these simulations make the decision-making process transparent and defensible [80]. They can demonstrate the potential impact on key metrics like safety, creator diversity, and the visibility of minority genres, facilitating a shared, evidence-based conversation about the desired direction of the platform. This approach aligns with the principles of the AI-Policy-Governance Nexus, which explains how regulatory pressure and AI integration reshape governance practices toward greater transparency and accountability [75]. The framework's emphasis on transparency directly mirrors Article 13 of the EU AI Act, which mandates that high-risk AI systems be designed to ensure their operations are sufficiently transparent for deployers to understand and use them appropriately [27] [83]. The requirement for developers to publicly summarize the content used for training general-purpose models is echoed in the `KO.IntegrityTrace`'s focus on provenance and disclosure [25].

The strategic implementation plan begins with the music domain and its variants, such as music videos and XR audio deployments in Phoenix, before expanding to other domains like film and gaming [132]. This phased approach is pragmatic, allowing for the stabilization of metrics and the refinement of models in a relatively contained environment. The core schema of the KOs is designed to be extensible. For example, fields like `modality` (audio, video, text), `duration`, and `interaction_type` (passive listening, interactive gameplay) can be added to adapt the framework to new contexts [22]. The application to XR environments introduces new complexities and opportunities. Public XR spaces, for instance, would require stringent safety gating for youth, leveraging the age-banded impact models from `KO.ImpactLog` [53]. The immersive nature of XR could amplify both harms (e.g., unsafe listening levels exceeding recommended maximums [46] [48]) and benefits (e.g., therapeutic applications for infants [47]). The cross-cultural dimension, while not a top priority for Phase 1, is acknowledged as a necessary evolution. The framework must eventually account for culture-specific norms regarding acceptable lyrical content or taboo topics [173]. This will involve developing scaling functions for `H` and `B` that reflect local values while adhering to global minimums for physical safety, such as a zero-tolerance policy for explicit incitement to violence [30]. The governance process for tuning these cultural thresholds must be transparent and auditable to prevent any single cultural group from gaining undue influence or "stakeholder privilege" [22].

Synthesis and Strategic Recommendations

This research agenda represents a pioneering and highly structured effort to embed empirical fairness into AI-driven entertainment systems. By treating fairness as a measurable object through the lens of Cybernet Knowledge Objects, the project moves decisively from abstract ethical principles to a concrete, testable, and implementable framework. The strategic prioritization of Topics 1 and 2—the development of `KO.IntegrityTrace` and `KO.ImpactLog`—is logical and foundational. These KOs establish the core inputs—the integrity and societal impact of creative works—that are necessary for making informed decisions about AI assistance and content routing. The simultaneous focus on Topics 4 and 7—bias audits and attribution robustness—is what elevates this initiative from a simple scoring mechanism to a truly responsible governance system. It proactively acknowledges that the pursuit of safety can create new forms of injustice and that attributing complex social phenomena to single works is a task of immense difficulty and responsibility.

The framework's strength lies in its dual focus on internal technical rigor and external policy transparency. The production of detailed JSON schemas and mathematical formulas ensures the system can be built and integrated into platforms like Cybernet/Reality.os. Concurrently, the creation of governance frameworks and impact simulations makes the system's logic and trade-offs legible to regulators, community partners in Phoenix, and the public at large. This dual-layered communication strategy is essential for building trust and fostering a collaborative approach to AI governance. The framework's design is forward-thinking, with a clear roadmap for evolving from a music-centric model to a cross-domain capability for film and gaming, and eventually to a globally adaptable system sensitive to cultural nuances. Its alignment with emerging regulatory paradigms, most notably the EU AI Act, positions it not just as a corporate initiative but as a potential model for responsible AI in creative industries.

To successfully execute this ambitious agenda, several actionable recommendations emerge from the analysis:

First, the initial phase of empirical validation should be laser-focused on developing and ground-truthing the constituent signals for the `Integrity_score`. This involves partnering with academic experts in Music Information Retrieval to refine plagiarism detection models like MelodySim and TruMuzic [1] [26]. Pilot studies should be conducted to establish a reliable methodology for measuring `marketing_truthfulness`, likely involving a hybrid of automated NLP and human oversight. Establishing expert-annotated datasets for these signals is a prerequisite for training accurate models.

Second, the research team must immediately begin developing a dedicated toolkit for causal inference to address the profound challenge of incident attribution. This toolkit should encompass a suite of statistical methods, such as difference-in-differences and propensity score matching, and the infrastructure to collect the necessary data, including exposure metrics, pre-intervention baselines, and control group data. The goal is to produce probabilistic attributions with confidence intervals, not deterministic verdicts, thereby instilling a culture of attributional humility into the system's core logic [97] [98].

Third, the conceptualization and design of the Minority Genre Protection Index (MGPI) should commence early in the project lifecycle. Defining what constitutes a "low-mainstream but culturally important genre" and brainstorming potential compensation mechanisms will ensure this critical bias-mitigation feature is a core, thoughtfully designed component of the system, rather than an afterthought. Historical precedents, such as royalty funds created to support artists displaced by technology, offer valuable models for designing these compensation schemes [43].

Finally, the dual-layer documentation strategy must be maintained throughout the project. The internal technical specifications should be treated as a living document, providing engineers with the precise formulas, data schemas, and logging requirements needed for implementation. The external governance and simulation documents should be regularly updated to reflect the latest findings and engage stakeholders in an ongoing dialogue about the framework's evolution. By adhering to this disciplined, empirically grounded, and ethically aware approach, this research has the potential to create a durable and equitable model for AI in entertainment that can serve as a benchmark for the industry.

Reference

1. MelodySim: Measuring Melody-aware Music Similarity for ... https://arxiv.org/html/2505.20979v1
2. Music fans believe AI companies training models with existing ... https://completemusicupdate.com/music-fans-believe-ai-companies-training-models-with-existing-recordings-must-get-consent-says-ifpi-study/
3. Regulating AI in the financial sector: recent developments ... https://www.bis.org/fsi/publ/insights63.pdf
4. IEEE P7001: A Proposed Standard on Transparency - PMC https://pmc.ncbi.nlm.nih.gov/articles/PMC8351056/
5. arXiv:2503.18814v1 [cs.AI] 24 Mar 2025 https://arxiv.org/pdf/2503.18814?
6. Law & Compliance in AI Security & Data Protection https://www.edpb.europa.eu/system/files/2025-06/spe-training-on-ai-and-data-protection-legal_en.pdf
7. The Role of Explainable AI in the Research Field of AI Ethics https://dl.acm.org/doi/full/10.1145/3599974
8. Heritability of Childhood Music Engagement and ... https://pmc.ncbi.nlm.nih.gov/articles/PMC10159991/
9. Understanding music and aging through the lens of ... https://www.sciencedirect.com/science/article/pii/S0149763424002379
10. Analysis of music-exposure interventions for impacting ... https://www.researchgate.net/journal/Current-Psychology-1936-4733/publication/368641653_Analysis_of_music-exposure_interventions_for_impacting_prosocial_behaviour_via_behaviour_change_techniques_and_mechanisms_of_action_a_rapid_review/links/63f1932419130a1a4a8e016b/Analysis-of-music-exposure-interventions-for-impacting-prosocial-behaviour-via-behaviour-change-techniques-and-mechanisms-of-action-a-rapid-review.pdf
11. Longitudinal Effects of Continuous Music Training on ... https://nyaspubs.onlinelibrary.wiley.com/doi/10.1111/nyas.70086
12. The Impact of Music on Human Development and Well-Being https://core.ac.uk/download/388547156.pdf
13. Mental health and music engagement: review, framework ... https://www.nature.com/articles/s41398-021-01483-8

14. Acceptability and effects on mental health of a music-based ... https://public-pages-files-2025.frontiersin.org/journals/virtual-reality/articles/10.3389/frvir.2025.1608416/pdf
15. Music and Stress https://www.researchgate.net/publication/225723923_Music_and_Stress
16. Mapping EEG Metrics to Human Affective and Cognitive Models https://pmc.ncbi.nlm.nih.gov/articles/PMC12649996/
17. The Relative Accuracy of Different Methods for Measuring ... https://onlinelibrary.wiley.com/doi/10.1002/brb3.70764
18. Developing An Effective Monitoring System Using Sensors ... https://research.nottingham.edu.cn/files/1204724292/Lingling_PhD_Thesis_Final.pdf
19. The neuroscience of sadness: A multidisciplinary synthesis ... https://www.sciencedirect.com/science/article/pii/S0149763418306146
20. Platforms, risk perceptions, and reporting: the impact of illicit ... https://pmc.ncbi.nlm.nih.gov/articles/PMC12495684/
21. understanding the individualized impacts of Internet use ... https://pmc.ncbi.nlm.nih.gov/articles/PMC11083903/
22. publishers' relationships with messaging platforms and the ... https://theses.hal.science/tel-04164939v1/file/103145_CARPES_DA_SILVA_2022_archivageb.pdf
23. Fairness in Music Recommender Systems: A Stakeholder ... https://pmc.ncbi.nlm.nih.gov/articles/PMC9353048/
24. Fighting Mainstream Bias in Recommender Systems via Local ... https://people.engr.tamu.edu/caverlee/pubs/Ziwei_WSDM_2022.pdf
25. IFPI looks at a decade of digital transformation in the music ... https://www.wipo.int/en/web/wipo-magazine/articles/ifpi-looks-at-a-decade-of-digital-transformation-in-the-music-industry-73661
26. TruMuzic: A Deep Learning and Data Provenance-Based ... https://www.mdpi.com/2076-3417/13/16/9425
27. AI Act Service Desk - Article 13: Transparency and provision of ... https://ai-act-service-desk.ec.europa.eu/en/ai-act/article-13
28. A Comprehensive Review on Harnessing Large Language ... https://arxiv.org/html/2507.21117v2
29. Opinion Paper: “So what if ChatGPT wrote it? ... https://www.sciencedirect.com/science/article/pii/S0268401223000233
30. Investigating how collective action mitigates the harm of ... https://www.sciencedirect.com/science/article/pii/S1471772725000314

31. A Collaborative, Human-Centred Taxonomy of AI, ... https://arxiv.org/html/2407.01294v1
32. Five dimensions for anticipating and repairing algorithmic harm https://journals.sagepub.com/doi/10.1177/20539517231211553
33. Bias in algorithms – Artificial intelligence and discrimination https://fra.europa.eu/sites/default/files/fra_uploads/fra-2022-bias-in-algorithms_en.pdf
34. (PDF) The impact of song-specific age and affective ... https://www.researchgate.net/publication/281139722_The_impact_of_song-specific_age_and_affective_qualities_of_popular_songs_on_music-evoked_autobiographical_memories_MEAMs
35. Datasets of Smartphone Modalities for Depression Assessment https://pmc.ncbi.nlm.nih.gov/articles/PMC12710863/
36. Poster Abstracts - 2023 - Early Intervention in Psychiatry https://onlinelibrary.wiley.com/doi/10.1111/eip.13409
37. Evaluating the Impact of Mindfulness Meditation to Reduce ... https://escholarship.org/content/qt9j815241/qt9j815241_noSplash_23739f200ae6171eb759abe82b26df02.pdf
38. Media violence and children: A complete guide for parents ... https://psycnet.apa.org/record/2004-00122-000
39. Developmental-Behavioral Pediatrics, Fifth Edition 2023 https://www.scribd.com/document/722649889/Developmental-Behavioral-Pediatrics-Fifth-Edition-2023
40. ALGORITHMS, ADDICTION, AND ADOLESCENT MENTAL ... https://www.cambridge.org/core/journals/american-journal-of-law-and-medicine/article/algorithms-addiction-and-adolescent-mental-health-an-interdisciplinary-study-to-inform-statelevel-policy-action-to-protect-youth-from-the-dangers-of-social-media/EC9754B533553BDD56827CD9E34DFC25
41. How TikTok Users Legitimate the Algorithmic Sensemaking https://journals.sagepub.com/doi/10.1177/20563051231224114
42. (PDF) AI and IP Rights Enforcement https://www.researchgate.net/publication/367255093_AI_and_IP_Rights_Enforcement
43. Protecting Human Creativity in AI-Generated Music with the ... https://academic.oup.com/grurint/article/73/12/1137/7832810
44. SoundSignature: What Type of Music Do You Like? https://arxiv.org/pdf/2410.03375?
45. Computer Science Jan 2023 http://arxiv.org/list/cs/2023-01?skip=2200&show=1000
46. Prevalence and global estimates of unsafe listening ... https://gh.bmj.com/content/7/11/e010501

47. Music Therapy in the Medical Care of Infants - Oxford Academic https://academic.oup.com/edited-volume/34473/chapter/292487651
48. (PDF) Decoding and spatial mapping of acoustic noise in ... https://www.researchgate.net/publication/395130611_Decoding_and_spatial_mapping_of_acoustic_noise_in_the_neonatal_intensive_care_unit
49. FluidTrack: Investigating Child-Parent Collaborative ... https://dl.acm.org/doi/10.1145/3706598.3713878
50. Measurement Method Options to Investigate Digital Screen ... https://pmc.ncbi.nlm.nih.gov/articles/PMC11275073/
51. Screen Time in Early Childhood Education https://documents1.worldbank.org/curated/en/099931007302544113/pdf/IDU-baaf52dd-348a-4f1c-a226-0f8d86c86094.pdf
52. Infants' Perception of Auditory Patterns (Chapter 8) https://www.cambridge.org/core/books/cambridge-handbook-of-infant-development/infants-perception-of-auditory-patterns/B2FB768232F30C9C0E09F8A230F86D2B
53. A Moving Metaverse: QoE challenges and standards ... https://www.sciencedirect.com/science/article/pii/S156849462400351X
54. Aligning XR Research with Autistic Priorities and Lived ... https://www.researchgate.net/publication/396865573_Aligning_XR_Research_with_Autistic_Priorities_and_Lived_Experiences_Insights_from_the_Project_PHoENIX_Study
55. Creating an Immersive XR Learning Experience https://www.mdpi.com/2079-9292/11/21/3547
56. 2. Architecture Requirements — CNTT-CNTT documentation https://cntt.readthedocs.io/en/stable-baraque/ref_arch/kubernetes/chapters/chapter02.html
57. OSA-Express Implementation Guide https://www.redbooks.ibm.com/redbooks/pdfs/sg245948.pdf
58. A systematic review of QoS enhancement techniques in ... https://dl.acm.org/doi/10.1016/j.compeleceng.2025.110550
59. DevNet Associate - API Architectural Styles | PDF https://www.scribd.com/document/775402917/DevNet-Associate-API-Architectural-Styles
60. Impacts of AI on Music Consumption and Fairness https://journals.sagepub.com/doi/10.1177/27523543241269047
61. Who Gets Heard? Rethinking Fairness in AI for Music ... https://arxiv.org/html/2511.05953v1

62. (PDF) Impacts of AI on Music Consumption and Fairness https://www.researchgate.net/publication/384133692_Impacts_of_AI_on_Music_Consumption_and_Fairness
63. An analysis of artificial intelligence automation in digital ... https://www.frontiersin.org/journals/artificial-intelligence/articles/10.3389/frai.2024.1515716/full
64. A survey on fairness-aware recommender systems https://www.sciencedirect.com/science/article/pii/S1566253523002221
65. ISO/IEC 23894:2023 - AI — Guidance on risk management https://www.iso.org/standard/77304.html
66. Fairness in Music Recommender Systems: A Stakeholder ... https://www.frontiersin.org/journals/big-data/articles/10.3389/fdata.2022.913608/full
67. A Survey on Trustworthy Recommender Systems https://dl.acm.org/doi/full/10.1145/3652891
68. A Deep Dive into Fairness, Bias, Threats, and Privacy in ... https://arxiv.org/pdf/2409.12651
69. The role of recommendation algorithms in the formation ... https://www.sciencedirect.com/science/article/pii/S0306457325001840
70. (PDF) Addressing the Fairness Issue of Large Music Models https://www.researchgate.net/publication/395723280_Addressing_the_Fairness_Issue_of_Large_Music_Models_A_Blockchain_Approach
71. assessing the FAIRness of data in social science research ... https://journals.sagepub.com/doi/10.1177/17470161241257575
72. Bias recognition and mitigation strategies in artificial ... https://pmc.ncbi.nlm.nih.gov/articles/PMC11897215/
73. AI-Based Risk Models for High-Power Individuals https://www.linkedin.com/pulse/ai-based-risk-models-high-power-individuals-andre-mcame
74. 10 AI dangers and risks and how to manage them https://www.ibm.com/think/insights/10-ai-dangers-and-risks-and-how-to-manage-them
75. The AI-policy-governance nexus: How regulation and ... https://www.sciencedirect.com/science/article/pii/S0160791X25003070
76. Cybernetic governance: implications of technology ... https://link.springer.com/article/10.1007/s10676-024-09763-9
77. It's just distributed computing: Rethinking AI governance https://www.sciencedirect.com/science/article/pii/S030859612500014X

78. Governance of Generative AI | Policy and Society https://academic.oup.com/policyandsociety/article/44/1/1/7997395
79. AI as a constituted system: accountability lessons from an ... https://www.cambridge.org/core/journals/data-and-policy/article/ai-as-a-constituted-system-accountability-lessons-from-an-llm-experiment/AA89E6ABD5189CBFF02FFECF4ED4F750
80. What is an AI governance framework? | AIGF explained https://talbotwest.com/services/ai-governance/what-is-an-ai-governance-framework
81. A Unified Framework for Cybernetic Orders in Self- ... https://www.linkedin.com/pulse/from-control-emergence-unified-framework-cybernetic-orders-kullok-gvnlf
82. Artificial intelligence, complexity, and systemic resilience in ... https://pmc.ncbi.nlm.nih.gov/articles/PMC12171231/
83. THE TRANSPARENCY MANDATE OF ARTICLE 13 IN THE ... https://www.linkedin.com/pulse/transparency-mandate-article-13-eu-ai-act-series-3-aumirah-1ovuc
84. Report claims' majority' of young people listen to AI music https://musically.com/2026/01/12/report-claims-majority-of-young-people-listen-to-ai-music/
85. Understanding AI Copyright Infringement Issues https://www.tiktok.com/@topmusicattorney/video/7522521347021458702
86. MelodySim: Measuring Melody-aware Music Similarity for ... https://arxiv.org/abs/2505.20979
87. Broadband networks of the future (EN) https://www.oecd.org/content/dam/oecd/en/publications/reports/2022/07/broadband-networks-of-the-future_cb6dd626/755e2d0c-en.pdf
88. Energy and AI - Microsoft .NET https://iea.blob.core.windows.net/assets/40a4db21-2225-42f0-8a07-addcc2ea86b3/EnergyandAI.pdf
89. Digital Intelligence Innovation https://chinamobileltd.com/en/esg/sd/2023/04.pdf
90. Speakers - AI for Good - ITU https://aiforgood.itu.int/summit25/speakers/
91. GUIDE TO THE MAN IN THE MIDDLE ATTACK (MITM) https://www.linkedin.com/pulse/guide-man-middle-attack-mitm-learn-what-how-defend-against-mamica-vrwjf
92. CyberSecurity Handbook and Reference Guide - Vol3 - 2018 https://pdfcoffee.com/cybersecurity-handbook-and-reference-guide-vol3-2018-pdf-free.html
93. 333333 23135851162 the 13151942776 of 12997637966 ftp://ftp.cs.princeton.edu/pub/cs226/autocomplete/words-333333.txt
94. (PDF) An analysis of artificial intelligence automation in ... https://www.researchgate.net/publication/

388531313_An_analysis_of_artificial_intelligence_automation_in_digital_music_streaming_platforms_for_improving_consumer_subscription_responses_a_review

95. Understanding the Dynamics in Deploying AI-Based ... https://dl.acm.org/doi/full/10.1145/3706598.3713532
96. Fairness in Recommender Systems: Evaluation ... https://dl.acm.org/doi/10.1145/3604558
97. Auditing the Ethical Logic of Generative AI Models https://arxiv.org/pdf/2504.17544
98. Algorithm Assurance: Auditing Applications of Artificial ... https://link.springer.com/chapter/10.1007/978-3-031-11089-4_7
99. AI alignment: Assessing the global impact of recommender ... https://www.sciencedirect.com/science/article/pii/S0016328724000661
100. (PDF) Algorithmic Fairness: A Tolerance Perspective https://www.researchgate.net/publication/380607650_Algorithmic_Fairness_A_Tolerance_Perspective
101. AI Act | Shaping Europe's digital future - European Union https://digital-strategy.ec.europa.eu/en/policies/regulatory-framework-ai
102. EU AI Act: first regulation on artificial intelligence | Topics https://www.europarl.europa.eu/topics/en/article/20230601STO93804/eu-ai-act-first-regulation-on-artificial-intelligence
103. arXiv:2412.01957v2 [cs.AI] 23 Jan 2025 https://arxiv.org/pdf/2412.01957
104. THE FUTURE OF AI AND AI GOVERNANCE IN AN ... https://www.itu.int/en/ITU-T/Workshops-and-Seminars/2023/0724/Documents/SayeghV2.pdf
105. AI creativity and legal protection for AI-generated works in ... https://www.sciencedirect.com/science/article/pii/S2666188825003156
106. rblx-20240630 https://www.sec.gov/Archives/edgar/data/1315098/000131509824000166/rblx-20240630.htm
107. AI, Machine Learning & Big Data 2024 | PDF https://www.scribd.com/document/751417190/AI-Machine-Learning-Big-Data-2024
108. Artificial Intelligence and the Rights of the Child https://publications.jrc.ec.europa.eu/repository/bitstream/JRC127564/JRC127564_01.pdf
109. Annual Report 2024 of Luxshare Precision Industry Co., Ltd. http://notice.10jqka.com.cn/api/pdf/7d48cafdc66068e6.pdf
110. Edge AI for Smart Cities: Foundations, Challenges, and ... https://www.researchgate.net/publication/398743598_Edge_AI_for_Smart_Cities_Foundations_Challenges_and_Opportunities
111. Computer Applications - Springer Link https://link.springer.com/content/pdf/10.1007/978-981-97-9674-8.pdf

112. Fairness Testing: A Comprehensive Survey and Analysis of ... https://dl.acm.org/doi/10.1145/3652155
113. A Stereotype-Based Framework for Addressing Trade-Offs https://www.researchgate.net/publication/396090778_Time-Evolving_Fairness_and_Accuracy_in_Recommendation_Systems_A_Stereotype-Based_Framework_for_Addressing_Trade-Offs
114. Applications of Artificial Intelligence in Music: A Review https://www.sciopen.com/article/10.26599/IJCS.2025.9100011
115. Algorithmic Fairness: A Tolerance Perspective https://arxiv.org/pdf/2405.09543?
116. The impact of algorithm awareness on the acceptance ... https://www.sciencedirect.com/science/article/pii/S0001691825006961
117. The Impact of AI-Personalized Recommendations on ... https://www.mdpi.com/0718-1876/20/1/21
118. Long-Term Fairness Inquiries and Pursuits in Machine Learn- ing https://openreview.net/notes/edits/attachment?id=VGUl9xykgl&name=pdf
119. Can Online Music Platforms Be Fair? An Interdisciplinary ... https://link.springer.com/article/10.1007/s40319-023-01420-w
120. A Green Paper on AI, Data Governance, and Metadata Policies ... https://zenodo.org/records/17229218/files/green-paper-v07.pdf?download=1
121. RecSys 2025 - Accepted Contributions https://recsys.acm.org/recsys25/accepted-contributions/
122. A Synthesis of Four Studies Conducted with Adult Learners https://www.researchgate.net/publication/393475892_A_cybernetic_guide_to_implementing_AI_for_collaborative_learning_A_synthesis_of_four_studies_conducted_with_adult_learners
123. A Survey of Cybersecurity Ontologies, Semantic Log ... https://arxiv.org/pdf/2510.16610
124. Ontology-based automated knowledge identification of ... https://www.tandfonline.com/doi/full/10.1080/13467581.2025.2467241
125. OECD AI Capability Indicators Technical Report https://www.oecd.org/content/dam/oecd/en/publications/reports/2025/11/oecd-ai-capability-indicators-technical-report_d3762d1a/9cdb3dd1-en.pdf
126. Driving Datafication (Part II) - International Economic Law ... https://www.cambridge.org/core/books/international-economic-law-in-the-era-of-datafication/driving-datafication/7CE07D0DF50163427802C26320BC3350
127. (PDF) The Influence of Media on Children's Psychological ... https://www.academia.edu/129054642/

The_Influence_of_Media_on_Children_s_Psychological_Development_and_Consumer_Behavior

128. Proceedings of the 2025 Conference on Empirical Methods ... https://aclanthology.org/volumes/2025.emnlp-industry/
129. Arxiv今日论文| 2025-12-24 http://lonepatient.top/2025/12/24/arxiv_papers_2025-12-24
130. An Empirical Study of Testing Machine Learning in the Wild https://dl.acm.org/doi/10.1145/3680463
131. Algorithms, Volume 18, Issue 12 (December 2025) https://www.mdpi.com/1999-4893/18/12
132. bing.txt ftp://ftp.cs.princeton.edu/pub/cs226/autocomplete/bing.txt
133. AI Watch Defining Artificial Intelligence 2.0 https://publications.jrc.ec.europa.eu/repository/bitstream/JRC126426/jrc126426_ai_watch_defining_artificial_intelligence_2.0_final_29-10-2021.pdf
134. A Scoping Review of Explainable AI Frameworks https://dl.acm.org/doi/full/10.1145/3769678
135. Thinking Beyond Tokens: From Brain-Inspired Intelligence ... https://arxiv.org/html/2507.00951v1
136. Artificial Intelligence Index Report 2025 https://hai-production.s3.amazonaws.com/files/hai_ai_index_report_2025.pdf
137. Mapping Generative AI rules and liability scenarios in the ... https://www.cambridge.org/core/journals/cambridge-forum-on-ai-law-and-governance/article/mapping-generative-ai-rules-and-liability-scenarios-in-the-ai-act-and-in-the-proposed-eu-liability-rules-for-ai-liability/CE2454E668DBD130E054996C6A984A31
138. Arxiv今日论文| 2025-12-17 - 闲记算法 http://lonepatient.top/2025/12/17/arxiv_papers_2025-12-17
139. 人工智能2025_9_9 http://www.arxivdaily.com/thread/71411
140. Track: Poster Session 1 - CVPR 2026 https://cvpr.thecvf.com/virtual/2025/session/35265
141. Technologies, Volume 13, Issue 8 (August 2025) – 66 articles https://www.mdpi.com/2227-7080/13/8
142. Memory Distortions Induced by Reality-Altering Media https://theses.hal.science/tel-05262805v1/file/137755_BONNAIL_2025_archivage.pdf
143. An analysis of artificial intelligence automation in digital music ... https://pmc.ncbi.nlm.nih.gov/articles/PMC11774895/
144. Comparing and modeling the use of online recommender ... https://www.sciencedirect.com/science/article/pii/S2451958824000824

145. Spotify tightens AI voice impersonation rules, introduces ... https://www.linkedin.com/posts/soundraw_spotifys-new-ai-protections-what-it-means-activity-7386259275817443328-Ueh4

146. ISO 42001, NIST AI RMF, and the EU AI Act https://www.linkedin.com/pulse/managing-ai-risks-iso-42001-nist-rmf-eu-act-your-faqs-oliver-qlcee

147. Music industry bodies welcome latest approval of EU's AI Act https://musically.com/2024/03/13/music-industry-bodies-welcome-latest-approval-of-eus-ai-act/

148. (PDF) A survey on popularity bias in recommender systems https://www.researchgate.net/publication/381913966_A_survey_on_popularity_bias_in_recommender_systems

149. (PDF) AI-DRIVEN DATA ENGINEERING TRANSFORMING ... https://www.researchgate.net/publication/386017582_AI-DRIVEN_DATA_ENGINEERING_TRANSFORMING_BIG_DATA_INTO_ACTIONABLE_INSIGHT

150. Wireless and Satellite Systems https://link.springer.com/content/pdf/10.1007/978-3-031-86196-3.pdf

151. №. 10 https://zenodo.org/records/15643187/files/issue.pdf?download=1

152. List of All | PDF https://www.scribd.com/document/643019297/list-of-all-xlsx

153. Dicionario portugues https://www.academia.edu/32592435/Dicionario_portugues

154. The Future of Child Development in the AI Era. Cross ... https://www.researchgate.net/publication/380974142_The_Future_of_Child_Development_in_the_AI_Era_Cross-Disciplinary_Perspectives_Between_AI_and_Child_Development_Experts

155. european commission - EUR-Lex https://eur-lex.europa.eu/legal-content/EN/TXT/HTML/?uri=CELEX:52020SC0348

156. (PDF) The DAta Protection REgulation COmpliance Model https://www.researchgate.net/publication/335953715_The_DAta_Protection_REgulation_COmpliance_Model

157. User Modeling and User Profiling: A Comprehensive Survey https://arxiv.org/html/2402.09660v2

158. Generative AI: A systematic review using topic modelling ... https://www.sciencedirect.com/science/article/pii/S2543925124000020

159. Natural Language Processing and Chinese Computing https://link.springer.com/content/pdf/10.1007/978-981-97-9440-9.pdf

160. Proceedings of the 62nd Annual Meeting of the Association ... https://aclanthology.org/2024.acl-demos.pdf

161. ScaleCUA: Scaling Open-Source Computer Use Agents ... https://www.researchgate.net/publication/395649749_ScaleCUA_Scaling_Open-Source_Computer_Use_Agents_with_Cross-Platform_Data
162. Everything Playaforms - S Meira | PDF | Business Model https://www.scribd.com/document/842264860/Everything-Playaforms-S-Meira
163. (PDF) Convergent Media and Privacy https://www.academia.edu/89289698/Convergent_Media_and_Privacy
164. A Scoping Review of Explainable AI Frameworks https://dl.acm.org/doi/10.1145/3769678
165. (PDF) Risk Management in the Artificial Intelligence Act https://www.researchgate.net/publication/368371151_Risk_Management_in_the_Artificial_Intelligence_Act
166. Foundation of Affective Computing & Interaction https://arxiv.org/html/2506.15497v1
167. A matter of choice: People and possibilities in the age of AI https://hdr.undp.org/system/files/documents/global-report-document/hdr2025reporten.pdf
168. The European AI Act: Requirements for High-Risk AI Systems https://www.emergobyul.com/news/european-ai-act-requirements-high-risk-ai-systems
169. A Survey on the Fairness of Recommender Systems https://dl.acm.org/doi/10.1145/3547333
170. Fairness in Recommender Systems: Research Landscape ... https://www.researchgate.net/publication/360777268_Fairness_in_Recommender_Systems_Research_Landscape_and_Future_Directions
171. A survey on popularity bias in recommender systems https://link.springer.com/article/10.1007/s11257-024-09406-0
172. arXiv:2008.01194v1 [cs.IR] 16 Jul 2020 https://arxiv.org/pdf/2008.01194
173. Part V - CONTEXTUAL FACTORS AND VIOLENT BEHAVIOR https://www.cambridge.org/core/books/cambridge-handbook-of-violent-behavior-and-aggression/contextual-factors-and-violent-behavior/91D00CE17DB082A0B4C9855A70A7C48A
174. Web and Big Data - Springer Link https://link.springer.com/content/pdf/10.1007/978-981-96-0055-7.pdf
175. https://www.researchgate.net/profile/David_Hale2/p... https://www.researchgate.net/profile/David_Hale2/publication/319160334_Primary_Study_Publications_excluded_as_NOT_being_related_to_SE_SLR_studies/data/5995c64b458515017ea59b54/SLR-Excluded-primary-studies.rtf

176. 259 - AVIONICS, AEROSPACE AND DEFENSE ... https://de.scribd.com/document/142763797/259-AVIONICS-AEROSPACE-AND-DEFENSE-ACRONYMS-AND-ABBREVIATIONS-Januar-2011
177. Artificial Intelligence in Education: The Intersection of ... https://www.scribd.com/document/843066206/978-3-031-71232-6
178. The Ethics of Advanced AI Assistants https://www.academia.edu/122521342/The_Ethics_of_Advanced_AI_Assistants
179. Guide To Cybersecurity in Digital Transformation Trends ... https://www.scribd.com/document/757807636/Guide-to-Cybersecurity-in-Digital-Transformation-Trends-Methods-Technologies-Applications-and-Best-Practices-Dietmar-P-F-Moller-Z-Library
180. 信息检索/信息论/社会&信息网络/CS与社会学2025_8_29 http://www.arxivdaily.com/thread/71093
181. CARANA Country Study Scenario-Based Exercise https://resourcehub01.blob.core.windows.net/training-files/Training%20Materials/020%20STM-UNSO/Ver24/020-128_CARANA%20Country%20Study_Complete%20Package.pdf
182. P172259091d4bc0260b75f074e... https://documents1.worldbank.org/curated/en/099140301092229225/txt/P172259091d4bc0260b75f074e8c93537f8.txt
183. Apple Developer Program License Agreement https://developer.apple.com/support/downloads/terms/apple-developer-program/Apple-Developer-Program-License-Agreement-English.pdf
184. Apple Developer Program License Agreement https://developer.apple.com/support/downloads/terms/apple-developer-program/Apple-Developer-Program-License-Agreement-English-UK.pdf
185. AI Governance in the Context of the EU AI Act https://www.researchgate.net/publication/394461661_AI_Governance_in_the_Context_of_the_EU_AI_Act
186. AI Governance - A Consolidated Reference | PDF https://www.scribd.com/document/909381127/AI-Governance-a-Consolidated-Reference
187. Poster Session 5 - CVPR 2026 https://cvpr.thecvf.com/virtual/2025/session/35269
188. Special Topic： 3D Point Cloud Processing and Applications https://www.zte.com.cn/content/dam/zte-site/res-www-zte-com-cn/mediares/magazine/publication/com_en/pdf/en202304.pdf
189. Spotify reveals its latest measures to handle AI ... https://musically.com/2025/09/25/spotify-reveals-its-latest-measures-to-handle-ai-music/
190. Open Source Projects https://docs.daocloud.io/native/open/
191. Spotify's Approach to Multi-Metric A/B Testing Decisions https://hackernoon.com/spotifys-approach-to-multi-metric-ab-testing-decisions

192. Environmental Social Governance Report https://investor.apple.com/files/doc_downloads/2021/08/2021_Apple_ESG_Report.pdf
193. Report on Risk Assessment and Risk Mitigation Measures https://www.apple.com/legal/dsa/20232808_app-store_risk-assessment-report_non-confidential.pdf
194. Machine Learning https://arxiv.org/list/cs.LG/new